

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt,seaborn as sns
```

```
In [2]: train_df=pd.read_csv(r"C:\Users\Lenovo\OneDrive\Desktop\Data Sets\Mobile_Price_Classification_train_data.csv")
train_df
```

Out[2]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749
...	...	...	...	...	...	...	...	...	...	...	...	...	...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	17	644
996	997	609	0	1.8	1	0	0	13	0.9	186	...	2	1152
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	12	477
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	12	38
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	19	457

1000 rows × 21 columns



```
In [3]: test_df=pd.read_csv(r"C:\Users\Lenovo\OneDrive\Desktop\Data Sets\Mobile_Price_Classification_test_data.csv")
test_df
```

Out[3]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208
...	...	...	...	...	...	...	...	...	...	...	...	...
1995	794	1	0.5	1	0	1	2	0.8	106	6	...	1222
1996	1965	1	2.6	1	0	0	39	0.2	187	4	...	915
1997	1911	0	0.9	1	1	1	36	0.7	108	8	...	868
1998	1512	0	0.9	0	4	1	46	0.1	145	5	...	336
1999	510	1	2.0	1	5	1	45	0.9	168	6	...	483

2000 rows × 21 columns



In [4]: train\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     1000 non-null   int64
1   battery_power          1000 non-null   int64
2   blue                   1000 non-null   int64
3   clock_speed            1000 non-null   float64
4   dual_sim               1000 non-null   int64
5   fc                     1000 non-null   int64
6   four_g                 1000 non-null   int64
7   int_memory             1000 non-null   int64
8   m_dep                  1000 non-null   float64
9   mobile_wt              1000 non-null   int64
10  n_cores                 1000 non-null   int64
11  pc                      1000 non-null   int64
12  px_height               1000 non-null   int64
13  px_width                1000 non-null   int64
14  ram                     1000 non-null   int64
15  sc_h                    1000 non-null   int64
16  sc_w                    1000 non-null   int64
17  talk_time               1000 non-null   int64
18  three_g                 1000 non-null   int64
19  touch_screen            1000 non-null   int64
20  wifi                    1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

In [5]: test\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   battery_power          2000 non-null   int64
1   blue                   2000 non-null   int64
2   clock_speed            2000 non-null   float64
3   dual_sim               2000 non-null   int64
4   fc                     2000 non-null   int64
5   four_g                 2000 non-null   int64
6   int_memory             2000 non-null   int64
7   m_dep                  2000 non-null   float64
8   mobile_wt              2000 non-null   int64
9   n_cores                 2000 non-null   int64
10  pc                      2000 non-null   int64
11  px_height               2000 non-null   int64
12  px_width                2000 non-null   int64
13  ram                     2000 non-null   int64
14  sc_h                    2000 non-null   int64
15  sc_w                    2000 non-null   int64
16  talk_time               2000 non-null   int64
17  three_g                 2000 non-null   int64
18  touch_screen            2000 non-null   int64
19  wifi                    2000 non-null   int64
20  price_range             2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.3 KB
```

```
In [7]: x=train_df.drop('wifi',axis=1)
        y=train_df['wifi']
```

```
In [9]: x=test_df.drop('wifi',axis=1)
        y=test_df['wifi']
```

```
In [10]: train_df['dual_sim'].value_counts()
```

```
Out[10]: dual_sim
1      517
0      483
Name: count, dtype: int64
```

```
In [11]: test_df['blue'].value_counts()
```

```
Out[11]: blue
0      1010
1       990
Name: count, dtype: int64
```

```
In [12]: T={"Home Owner":{"Yes":1,"No":0}}
        train_df=train_df.replace(T)
        print(train_df)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	1	1043	1	1.8	1	14	0	5
1	2	841	1	0.5	1	4	1	61
2	3	1807	1	2.8	0	1	0	27
3	4	1546	0	0.5	1	18	1	25
4	5	1434	0	1.4	0	11	1	49
..	...	...	...	...	...	..	...	...
995	996	1700	1	1.9	0	0	1	54
996	997	609	0	1.8	1	0	0	13
997	998	1185	0	1.4	0	1	1	8
998	999	1533	1	0.5	1	0	0	50
999	1000	1270	1	0.5	0	4	1	35

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w
0	0.1	193	...	16	226	1412	3476	12	7
1	0.8	191	...	12	746	857	3895	6	0
2	0.9	186	...	4	1270	1366	2396	17	10
3	0.5	96	...	20	295	1752	3893	10	0
4	0.5	108	...	18	749	810	1773	15	8
..	...	...	...	..	...	...	...	...	...
995	0.5	170	...	17	644	913	2121	14	8
996	0.9	186	...	2	1152	1632	1933	8	1
997	0.5	80	...	12	477	825	1223	5	0
998	0.4	171	...	12	38	832	2509	15	11
999	0.1	140	...	19	457	608	2828	9	2

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..	...	...	...	...
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

```
In [13]: T={"Home Owner":{"Yes":1,"No":0}}
train_df=train_df.replace(T)
print(train_df)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	1	1043	1	1.8	1	14	0	5	\
1	2	841	1	0.5	1	4	1	61	
2	3	1807	1	2.8	0	1	0	27	
3	4	1546	0	0.5	1	18	1	25	
4	5	1434	0	1.4	0	11	1	49	
..	...	...	...	...	...	..	...	...	
995	996	1700	1	1.9	0	0	1	54	
996	997	609	0	1.8	1	0	0	13	
997	998	1185	0	1.4	0	1	1	8	
998	999	1533	1	0.5	1	0	0	50	
999	1000	1270	1	0.5	0	4	1	35	

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	
0	0.1	193	...	16	226	1412	3476	12	7	\
1	0.8	191	...	12	746	857	3895	6	0	
2	0.9	186	...	4	1270	1366	2396	17	10	
3	0.5	96	...	20	295	1752	3893	10	0	
4	0.5	108	...	18	749	810	1773	15	8	
..	...	...	...	..	...	...	...	...	...	
995	0.5	170	...	17	644	913	2121	14	8	
996	0.9	186	...	2	1152	1632	1933	8	1	
997	0.5	80	...	12	477	825	1223	5	0	
998	0.4	171	...	12	38	832	2509	15	11	
999	0.1	140	...	19	457	608	2828	9	2	

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..	...	...	...	...
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

```
In [14]: T={"Home Owner":{"Yes":1,"No":0}}
test_df=test_df.replace(T)
print(test_df)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	842	0	2.2	0	1	0	7	\
1	1021	1	0.5	1	0	1	53	
2	563	1	0.5	1	2	1	41	
3	615	1	2.5	0	0	0	10	
4	1821	1	1.2	0	13	1	44	
...	...	...	...	...	...	...	...	
1995	794	1	0.5	1	0	1	2	
1996	1965	1	2.6	1	0	0	39	
1997	1911	0	0.9	1	1	1	36	
1998	1512	0	0.9	0	4	1	46	
1999	510	1	2.0	1	5	1	45	

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	
0	0.6	188	2	...	20	756	2549	9	7	\
1	0.7	136	3	...	905	1988	2631	17	3	
2	0.9	145	5	...	1263	1716	2603	11	2	
3	0.8	131	6	...	1216	1786	2769	16	8	
4	0.6	141	2	...	1208	1212	1411	8	2	
...	...	...	...	...	...	...	...	...	...	
1995	0.8	106	6	...	1222	1890	668	13	4	
1996	0.2	187	4	...	915	1965	2032	11	10	
1997	0.7	108	8	...	868	1632	3057	9	1	
1998	0.1	145	5	...	336	670	869	18	10	
1999	0.9	168	6	...	483	754	3919	19	4	

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...	...	...	...	...	...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

```
In [15]: x=train_df.drop('wifi',axis=1)
y=train_df['wifi']
```

```
In [16]: x=test_df.drop('wifi',axis=1)
y=test_df['wifi']
```

```
In [17]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

```
Out[17]: ((1400, 20), (600, 20))
```

```
In [18]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[18]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [19]: rf = RandomForestClassifier()
```

```
In [20]: params = {'max_depth': [2,3,5,10,20],
'min_samples_leaf': [5,10,20,50,100,200],
'n_estimators': [10,25,30,50,100,200]}
```

```
In [21]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rf,param_grid=params,cv = 2, scoring='accuracy')
grid_search.fit(x_train,y_train)
```

```
Out[21]: ▸ GridSearchCV
▸ estimator: RandomForestClassifier
▸ RandomForestClassifier
```

```
In [22]: grid_search.best_score_
```

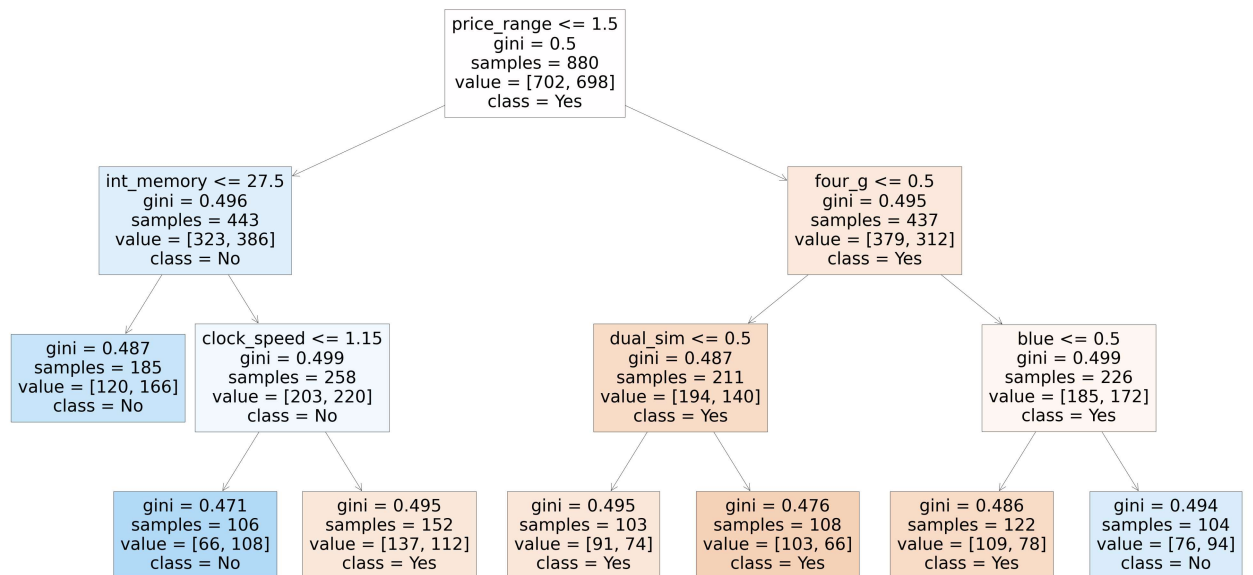
```
Out[22]: 0.5285714285714285
```

```
In [23]: rf_best = grid_search.best_estimator_
print(rf_best)
```

```
RandomForestClassifier(max_depth=3, min_samples_leaf=100, n_estimators=10)
```

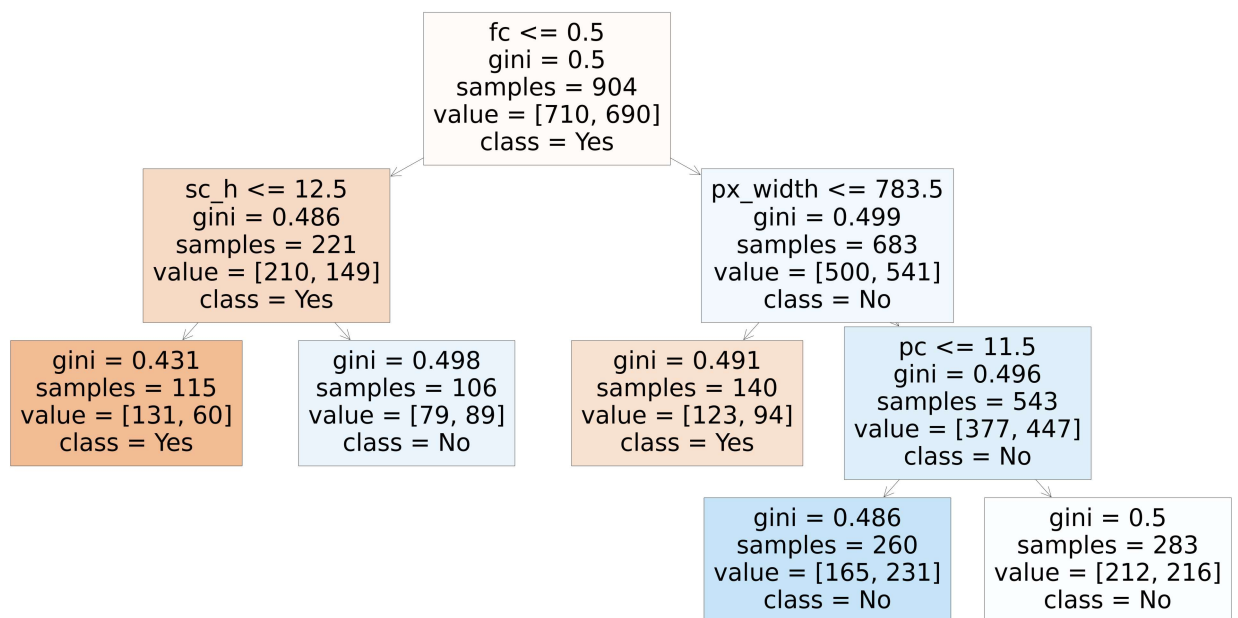
```
In [24]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5], feature_names = x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[24]: [Text(0.4230769230769231, 0.875, 'price_range <= 1.5\ngini = 0.5\nsamples = 880\nvalue = [702, 698]\nnclass = Yes'),
Text(0.15384615384615385, 0.625, 'int_memory <= 27.5\ngini = 0.496\nsamples = 443\nvalue = [323, 386]\nnclass = No'),
Text(0.07692307692307693, 0.375, 'gini = 0.487\nsamples = 185\nvalue = [120, 166]\nnclass = No'),
Text(0.23076923076923078, 0.375, 'clock_speed <= 1.15\ngini = 0.499\nsamples = 258\nvalue = [203, 220]\nnclass = No'),
Text(0.15384615384615385, 0.125, 'gini = 0.471\nsamples = 106\nvalue = [66, 108]\nnclass = No'),
Text(0.3076923076923077, 0.125, 'gini = 0.495\nsamples = 152\nvalue = [137, 112]\nnclass = Yes'),
Text(0.6923076923076923, 0.625, 'four_g <= 0.5\ngini = 0.495\nsamples = 437\nvalue = [379, 312]\nnclass = Yes'),
Text(0.5384615384615384, 0.375, 'dual_sim <= 0.5\ngini = 0.487\nsamples = 211\nvalue = [194, 140]\nnclass = Yes'),
Text(0.46153846153846156, 0.125, 'gini = 0.495\nsamples = 103\nvalue = [91, 74]\nnclass = Yes'),
Text(0.6153846153846154, 0.125, 'gini = 0.476\nsamples = 108\nvalue = [103, 66]\nnclass = Yes'),
Text(0.8461538461538461, 0.375, 'blue <= 0.5\ngini = 0.499\nsamples = 226\nvalue = [185, 172]\nnclass = Yes'),
Text(0.7692307692307693, 0.125, 'gini = 0.486\nsamples = 122\nvalue = [109, 78]\nnclass = Yes'),
Text(0.9230769230769231, 0.125, 'gini = 0.494\nsamples = 104\nvalue = [76, 94]\nnclass = No')]
```



```
In [25]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["Yes","No"],filled=True)
```

```
Out[25]: [Text(0.4444444444444444, 0.875, 'fc <= 0.5\ngini = 0.5\nsamples = 904\nvalue = [710, 690]\ncla
ss = Yes'),
Text(0.2222222222222222, 0.625, 'sc_h <= 12.5\ngini = 0.486\nsamples = 221\nvalue = [210, 149]
\nclass = Yes'),
Text(0.1111111111111111, 0.375, 'gini = 0.431\nsamples = 115\nvalue = [131, 60]\nnclass = Ye
s'),
Text(0.3333333333333333, 0.375, 'gini = 0.498\nsamples = 106\nvalue = [79, 89]\nnclass = No'),
Text(0.6666666666666666, 0.625, 'px_width <= 783.5\ngini = 0.499\nsamples = 683\nvalue = [500,
541]\nnclass = No'),
Text(0.5555555555555555, 0.375, 'gini = 0.491\nsamples = 140\nvalue = [123, 94]\nnclass = Ye
s'),
Text(0.7777777777777778, 0.375, 'pc <= 11.5\ngini = 0.496\nsamples = 543\nvalue = [377, 447]\n
class = No'),
Text(0.6666666666666666, 0.125, 'gini = 0.486\nsamples = 260\nvalue = [165, 231]\nnclass = N
o'),
Text(0.8888888888888888, 0.125, 'gini = 0.5\nsamples = 283\nvalue = [212, 216]\nnclass = No')]
```



```
In [26]: rf_best.feature_importances_
```

```
Out[26]: array([0.02266818, 0.01779351, 0.08494914, 0.00303683, 0.05861945,
0.05858221, 0.06054766, 0.00585816, 0.04014095, 0.
0.0543034 , 0.12675396, 0.12083748, 0.14386263, 0.0515548 ,
0.00618813, 0.05709013, 0.
, 0.02358964, 0.06362373])
```



```
In [27]: imp_df = pd.DataFrame({"Vername": x_train.columns, "Imp": rf_best.feature_importances_})
imp_df.sort_values(by="Imp", ascending=False)
```

Out[27]:

	Vername	Imp
13	ram	0.143863
11	px_height	0.126754
12	px_width	0.120837
2	clock_speed	0.084949
19	price_range	0.063624
6	int_memory	0.060548
4	fc	0.058619
5	four_g	0.058582
16	talk_time	0.057090
10	pc	0.054303
14	sc_h	0.051555
8	mobile_wt	0.040141
18	touch_screen	0.023590
0	battery_power	0.022668
1	blue	0.017794
15	sc_w	0.006188
7	m_dep	0.005858
3	dual_sim	0.003037
9	n_cores	0.000000
17	three_g	0.000000